

Max-Diversity Distributed Learning: Theory and Algorithms

Yong Liu, Jian Li, Weiping Wang

Abstract—We study the risk performance of distributed learning for the regularization empirical risk minimization with fast convergence rate, substantially improving the error analysis of the existing divide-and-conquer based distributed learning. An interesting theoretical finding is that the larger the diversity of each local estimate is, the tighter the risk bound is. This theoretical analysis motivates us to devise an effective max-diversity distributed learning algorithm (MDD). Experimental results show that MDD can outperform the existing divide-and-conquer methods but with a bit more time. Theoretical analysis and empirical results demonstrate that our proposed MDD is sound and effective.

Index Terms—Distributed Learning, Empirical Risk Minimization

I. INTRODUCTION

In recent years, the rapid expansion of data size brings a series of scientific challenges, such as algorithmic scalability and storage bottleneck [1], [2], [3]. Distributed learning based on a divide-and-conquer approach has popular used in various areas [4], [5], [6]. It breaks up a big problem into manageable pieces, operates learning algorithms on each pieces, and finally puts the individual solutions together to obtain a global output.

In this paper, we focus on the error analysis of the distributed learning for (regularization) empirical risk minimization. Given

$$\mathcal{S} = \{z_i = (\mathbf{x}_i, y_i)\}_{i=1}^N \in (\mathcal{Z} = \mathcal{X} \times \mathcal{Y})^N,$$

drawn identically and independently (i.i.d) from a probability distribution \mathbb{P} on \mathcal{Z} , the optimization problem of (regularization) empirical risk minimization can be stated as

$$\hat{f} = \arg \min_{f \in \mathcal{H}} \hat{R}(f) := \frac{1}{N} \sum_{j=1}^N \ell(f, z_j) + r(f) \quad (1)$$

where $\ell(f, z)$ is a loss function, $r(f)$ is a regularizer, and \mathcal{H} is a hypothesis space. This method has been well studied in learning theory, see e.g. [7], [8], [9], [10], [11]. In distributed learning, the data set \mathcal{S} is partitioned into m disjoint subsets

This work is supported in part by the National Key Research and Development Program of China (No.2016YFB1000604), the Science and Technology Project of Beijing (No.Z181100002718004), the National Natural Science Foundation of China (No.6173396, No.61673293, No.61602467) and the Excellent Talent Introduction of Institute of Information Engineering of CAS (Y7Z0111107).

Y. Liu, J. Li and W.P. Wang are with the Institute of Information Engineering, Chinese Academy of Sciences (e-mail: liuyong@iie.ac.cn).

$\{\mathcal{S}_i\}_{i=1}^m$, $|\mathcal{S}_i| = \frac{N}{m} =: n$. Then it assigns each \mathcal{S}_i to one machine or processor to get a local estimator \hat{f}_i :

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \hat{R}_i(f) := \frac{1}{|\mathcal{S}_i|} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f).$$

The finally global estimator \bar{f} is synthesized by

$$\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i.$$

Theoretical analysis of divide-and-conquer based distributed learning form a hot topic in machine learning [4], [2], [3], [12]. Under local strong convexity, smoothness and a reasonable set of other conditions, [4] showed that the mean-squared error decays as

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{1}{n^2} \right),$$

where f^* is the optimal hypothesis in the hypothesis space. Under some eigenfunction assumption, the error analysis for distributed kernel-based least squares was established in [2]: if m is not too large,

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\|f^*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N} \right),$$

where $\gamma(\lambda) = \sum_{j=1}^{\infty} \frac{\mu_j}{\lambda + \mu_j}$, μ_j is the eigenvalue of a Mercer kernel function. Without any eigenfunction assumption, an improved bound was derived for some $1 \leq p \leq \infty$ [3]:

$$\mathbb{E} [\|\bar{f} - f^*\|] = \mathcal{O} \left(\left(\frac{\gamma(\lambda)}{N} \right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N} \right)^{\frac{1}{2p}} \right).$$

There are two main contributions in this paper. First, under strongly convex and smooth, and a reasonable set of other conditions, we derive a risk bound of fast rate:

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_*}{n} + \frac{1}{n^2} - \Delta_{\bar{f}} \right), \quad (2)$$

where

$$\Delta_{\bar{f}} = \mathcal{O} \left(\frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2 \right)$$

is the diversity between all partition-based estimates, $R(f) = \mathbb{E}_z [\ell(f, z)] + r(f)$, and $H_* = \mathbb{E}_z [\ell(f_*, z)]$. When the minimal risk is small, the rate can reach

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{1}{n^2} - \Delta_{\bar{f}} \right).$$

Thus, if $m \leq \sqrt{N}$, the order of $R(\bar{f}) - R(f_*)$ is faster than $\mathcal{O}(\frac{1}{N} - \Delta_{\bar{f}})$. Note that if $\ell(f, z) + r(f)$ is L -Lipschitz continuous over f , the order of $R(\bar{f}) - R(f_*)$ is

$$\begin{aligned} R(\bar{f}) - R(f_*) &= \mathcal{O}(L\mathbb{E}[\|\bar{f} - f^*\|]) \\ &= \mathcal{O}\left(L\sqrt{\mathbb{E}[\|\bar{f} - f^*\|^2]}\right). \end{aligned}$$

Thus, the order of $R(\bar{f}) - R(f_*)$ in [2], [3], [4] at most $\mathcal{O}(\frac{1}{\sqrt{N}})$, which is much slower than that of our bound. Our second contribution is to design a novel max-diversity distributed learning algorithm. From Equation (2), we know that the larger the diversity $\Delta_{\bar{f}}$ is, the tighter the risk bound is. This interesting theoretical finding motivates us to devise a max-diversity distributed learning algorithm (MDD):

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \frac{1}{|S_i|} \sum_{z_j \in S_i} \ell(f, z_j) + r(f) - \gamma \|f - \bar{f}_i\|_{\mathcal{H}}, \quad (3)$$

where

$$\bar{f}_i = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{f}_j.$$

The last term of (3) is to make $\Delta_{\bar{f}}$ large. Experimental results show that our MDD can outperform the existing divide-and-conquer methods but with a bit more computational cost. As far as we know, the theoretical results w.r.t. diversity are given for a distributed setting has never given before.

The rest of this paper is organized as follows. We derive a risk bound of distributed learning with fast convergence rate in Section 2. In Section 3, we propose two novel algorithms based on the max-diversity of each local estimate in linear space and RKHS. In Section 4, we empirically study the performance of our MDD. We end in Section 5 with conclusion. All the proofs are given in the last part.

II. ERROR ANALYSIS OF DISTRIBUTED LEARNING

We will derive a sharper risk bound under some common assumptions in this section.

A. Assumptions

In the following, $\|\cdot\|_{\mathcal{H}}$ is denoted as the norm of the Hilbert space \mathcal{H} . The expected risk $R(f)$ and the optimal hypothesis f_* are denoted as

$$R(f) = \mathbb{E}_z[\ell(f, z)] + r(f) \text{ and } f_* = \arg \min_{f \in \mathcal{H}} R(f).$$

Assumption 1. The risk $R(f)$ is an η -strongly convex function, that is $\forall f, f' \in \mathcal{H}$,

$$\langle \nabla R(f), f - f' \rangle_{\mathcal{H}} + \frac{\eta}{2} \|f - f'\|_{\mathcal{H}}^2 \leq R(f) - R(f'), \quad (4)$$

or (another equivalent definition) $\forall f, f' \in \mathcal{H}, t \in [0, 1]$,

$$\begin{aligned} &R(tf + (1-t)f') \\ &\leq tR(f) + (1-t)R(f') - \frac{1}{2}\eta t(t-1)\|f - f'\|_{\mathcal{H}}^2. \end{aligned} \quad (5)$$

Assumption 2. The empirical risk $\hat{R}(f)$ is a convex function.

Assumption 3. The loss function $\ell(f, z)$ is τ -smooth with respect to the first variable f , that is $\forall f, f' \in \mathcal{H}$,

$$\|\nabla \ell(f, \cdot) - \nabla \ell(f', \cdot)\|_{\mathcal{H}} \leq \tau \|f - f'\|_{\mathcal{H}}. \quad (6)$$

Assumption 4. The regularizer $r(f)$ is a τ' -smooth function, that is $\forall f, f' \in \mathcal{H}$,

$$\|\nabla r(f) - \nabla r(f')\|_{\mathcal{H}} \leq \tau' \|f - f'\|_{\mathcal{H}}. \quad (7)$$

Assumption 5. The function $\nu(f, z) = \ell(f, z) + r(f)$ is L -Lipschitz continuous with respect to the first variable f , that is $\forall f, f' \in \mathcal{H}$,

$$\|\nu(f, \cdot) - \nu(f', \cdot)\|_{\mathcal{H}} \leq L \|f - f'\|_{\mathcal{H}}. \quad (8)$$

Assumptions 1, 2, 3, 4 and 5 allow us to model some popular losses, such as square loss and logistic loss, and some regularizer, such as $r(f) = \lambda \|f\|_{\mathcal{H}}^2$.

Assumption 6. We assume that the gradient at f_* is upper bounded by M , that is

$$\|\nabla \ell(f_*, \cdot)\|_{\mathcal{H}} \leq M.$$

Assumption 6 is also a common assumption, which is used in [13], [4].

B. Faster Rate of Distributed Learning

Let $\mathcal{N}(\mathcal{H}, \epsilon)$ be the ϵ -net of \mathcal{H} with minimal cardinality, and $C(\mathcal{H}, \epsilon)$ the covering number of $|\mathcal{N}(\mathcal{H}, \epsilon)|$

Theorem 1. For any $0 < \delta < 1$, $\epsilon \geq 0$, under **Assumptions 1, 2, 3, 4, 5 and 6**, and when

$$m \leq \frac{N\eta}{4\tilde{\tau} \log C(\mathcal{H}, \epsilon)}, \quad (9)$$

at least $1 - \delta$, we have

$$\begin{aligned} &R(\bar{f}) - R(f_*) \\ &\leq \frac{16\tilde{\tau} \log(4m/\delta)}{n^2\eta} + \frac{128\tau H_* \log(4m/\delta)}{n\eta} \\ &\quad + \frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} \\ &\quad + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta} - \Delta_{\bar{f}}, \end{aligned} \quad (10)$$

where $\Delta_{\bar{f}} = \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2$, $H_* = \mathbb{E}_z[\ell(f_*, z)]$ and $\tilde{\tau} = \tau + \tau'$.

From the above theorem, an interesting finding is that, when the larger the diversity of each local estimate is, the tighter the risk bound is. Furthermore, one can also see that when ϵ small enough,

$$\frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta}$$

will become non-dominating. To be specific, we have the following result:

Corollary 1. Under the same assumptions as Theorem 1, and setting $\epsilon = \frac{1}{n}$, when $m \leq \frac{N\eta}{4\tilde{\tau} \log C(\mathcal{H}, 1/n)}$, with high probability, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O}\left(\frac{H_* \log(m)}{n} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}}\right).$$

If the the minimal risk H_* is small, i.e., $H_* \leq \mathcal{O}(\frac{1}{n})$, the rate can reach

$$\mathcal{O}\left(\frac{\log(m)}{n^2} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}}\right).$$

As far as we know, the $\tilde{\mathcal{O}}(\frac{1}{n^2})$ -type of distributed risk bound for (regularization) empirical risk minimization has not been given before.

Remark 1. Note that if the hypothesis space \mathcal{H} is not very bad, the risk of the optimal hypothesis in \mathcal{H} is usually small. Thus, the assumption of the minimal risk H_* is small is a mild assumption.

In the next, we will consider two popular hypothesis spaces: linear and reproducing kernel Hilbert space (RKHS).

C. Linear Space

The linear hypothesis space we considered is defined as

$$\mathcal{H} = \left\{ f = \mathbf{w}^T \mathbf{x} \mid \mathbf{w} \in \mathbb{R}^d, \|\mathbf{w}\|_2 \leq B \right\}.$$

From [14], the cover number of linear hypothesis space can be bounded by

$$\log(C(\mathcal{H}, \epsilon)) \leq d \log\left(\frac{6B}{\epsilon}\right).$$

Thus, if we set $\epsilon = \frac{1}{n}$, from Corollary 1, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O}\left(\frac{H_* \log m}{n} + \frac{d \log n}{n^2} - \Delta_{\bar{f}}\right)$$

If the minimal risk $H_* \leq \mathcal{O}(\frac{d}{n})$, the rate can reach

$$\mathcal{O}\left(\frac{d \log(mn)}{n^2} - \Delta_{\bar{f}}\right) = \mathcal{O}\left(\frac{d \log N}{n^2} - \Delta_{\bar{f}}\right).$$

Therefore, if $m \leq \sqrt{\frac{N}{d \log N}}$, the order of risk bound can even faster than $\mathcal{O}(\frac{1}{N})$.

D. Reproducing Kernel Hilbert Space (RKHS)

The RKHS \mathcal{H}_K induced by the kernel function K is defined to be the closure of the linear span of the set of functions $\{K(\mathbf{x}, \cdot) : \mathbf{x} \in \mathcal{X}\}$ with the inner product satisfying

$$\langle K(\mathbf{x}, \cdot), f \rangle_K = f(\mathbf{x}), \forall \mathbf{x} \in \mathcal{X}, f \in \mathcal{H}_K.$$

The hypothesis space of RKHS in this paper is

$$\mathcal{H} := \{f \in \mathcal{H}_K : \|f\|_{\mathcal{H}_K} \leq B\}.$$

From [15], if K is the popular Gaussian kernel over $[0, 1]^d$:

$$K(\mathbf{x}, \mathbf{x}') = \exp\left\{-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{\sigma^2}\right\}, \mathbf{x}, \mathbf{x}' \in [0, 1]^d,$$

then for $0 \leq \epsilon \leq \frac{B}{2}$,

$$\log(C(\mathcal{H}, \epsilon)) = \mathcal{O}\left(\log^d\left(\frac{B}{\epsilon}\right)\right).$$

From Corollary 1, if we set $\epsilon = \frac{1}{n}$, and assume $H_* \leq \mathcal{O}(\frac{1}{n})$,

$$R(\bar{f}) - R(f_*) = \mathcal{O}\left(\frac{\log m}{n^2} + \frac{\log^d n}{n^2} - \Delta_{\bar{f}}\right)$$

Therefore, if $m \leq \min\left\{\sqrt{\frac{N}{d \log N}}, \sqrt{\frac{N}{\log^d n}}\right\}$, the order is faster than $\mathcal{O}(\frac{1}{N})$.

E. Comparison with Related Work

In this subsection, we compare our result with the most related work [2], [3], [4]. Under the smooth, strongly convex and other some assumptions, a distributed risk bound is given in [4]:

$$\mathbb{E}[\|\bar{f} - f_*\|^2] = \mathcal{O}\left(\frac{1}{N} + \frac{\log d}{n^2}\right).$$

Under some eigenfunction assumption, the error analysis for distributed regularized least squares were established in [2],

$$\mathbb{E}[\|\bar{f} - f^*\|^2] = \mathcal{O}\left(\|f_*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N}\right).$$

By removing the eigenfunction assumptions with a novel integral operator method of [2], a new bound was derived [3]:

$$\mathbb{E}[\|\bar{f} - f^*\|] = \mathcal{O}\left(\left(\frac{\gamma(\lambda)}{N}\right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N}\right)^{\frac{1}{2p}}\right).$$

Note that, if $\nu(f, z)$ is L -Lipschitz continuous over f , that is

$$\forall f, f' \in \mathcal{H}, z \in \mathcal{Z}, |\nu(f, z) - \nu(f', z)| \leq L\|f - f'\|,$$

we can obtain that

$$R(f) - R(f_*) \leq L \mathbb{E}[\|\bar{f} - f_*\|] \leq L \sqrt{\mathbb{E}[\|\bar{f} - f_*\|^2]}.$$

Thus, the order of [2], [3], [4] of $R(f) - R(f_*)$ is at most $\mathcal{O}\left(\frac{1}{\sqrt{N}}\right)$.

According to the subsections II-C and II-D, if m is not very large, and H_* is small, the order of this paper can even faster than $\mathcal{O}(\frac{1}{N})$, which is much faster than those of in the related work [4], [2], [3].

III. MAX-DISCREPANT DISTRIBUTED LEARNING (MDD)

In this section, we will propose two novel algorithms for linear space and RKHS, respectively. From corollary 1, we know that

$$R(f) - R(f_*) = \mathcal{O}\left(\frac{1}{n^2} - \frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2\right).$$

Thus, to obtain tighter bound, the diversity of each local estimate $\hat{f}_i, i = 1, \dots, m$, should be larger.

A. Linear Hypothesis Space

When \mathcal{H} is a linear Hypothesis space, we consider the following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 + \gamma \mathbf{w}^T \bar{\mathbf{w}}_{\setminus i}, \quad (11)$$

where $\bar{\mathbf{w}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{\mathbf{w}}_j$. Note that, if given $\bar{\mathbf{w}}_{\setminus i}$, $\hat{\mathbf{w}}_i$ has following closed form solution:

$$\hat{\mathbf{w}}_i = \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d}_{:= \mathbf{A}_i} \right)^{-1} \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}}_{:= \mathbf{b}_i} - \frac{\gamma \bar{\mathbf{w}}_{\setminus i}}{2} \right),$$

where $\mathbf{X}_{\mathcal{S}_i} = (\mathbf{x}_{t_1}, \mathbf{x}_{t_2}, \dots, \mathbf{x}_{t_n})$, $\mathbf{y}_{\mathcal{S}_i} = (y_{t_1}, y_{t_2}, \dots, y_{t_n})^T$, $z_{t_j} \in \mathcal{S}_i$, $j = 1, \dots, n$. In the next, we will give an iterative algorithm to solve the optimization problem (11). In each iteration, we should compute $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$ that needs $\mathcal{O}(d^2)$ if given \mathbf{A}_i^{-1} , which is computational intensive. Fortunately, from Lemma 3 (see in the last part), the $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$ can be computed by

$$\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i} = \left(\bar{\mathbf{w}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i,$$

where $a./c = (a/c_1, \dots, a/c_d)^T$, which only needs $\mathcal{O}(d)$.

The Max-Discrepant Distributed Learning algorithm for linear space is given in Algorithm 1. Compared with the traditional divide-and-conquer method, our MDD for linear space only need add $\mathcal{O}(d)$ in each iteration for each worker node.

Algorithm 1 Max-Discrepant Distributed Learning for Linear Space (MDD-LS)

- 1: **Input:** $\lambda, \gamma, \mathbf{X}, m, \zeta$.
 - 2: **For each worker node i :** $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$, and push $\hat{\mathbf{w}}_i^0$ to the server node.
 $// \mathbf{A}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d, \mathbf{b}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}$.
 - 3: **For server node:** $\bar{\mathbf{w}}^0 = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^0$, $\bar{\mathbf{w}}_{\setminus i}^0 = \frac{m \bar{\mathbf{w}}^0 - \hat{\mathbf{w}}_i^0}{m-1}$.
 - 4: **for** $t = 1, 2, \dots$ **do**
 - 5: **For each worker node i :**
 Pull $\bar{\mathbf{w}}_{\setminus i}^{t-1}$ from server node.
 $\mathbf{d}_i^t = \left(\left(\bar{\mathbf{w}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) ./ \mathbf{b}_i, \hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$.
 - 6: Push $\hat{\mathbf{w}}_i^t$ to the server node.
 - 7: **For server node:**
 $\bar{\mathbf{w}}^t = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$
 if $\|\bar{\mathbf{w}}^t - \bar{\mathbf{w}}^{t-1}\| \leq \zeta$ **end for**
 else $\bar{\mathbf{w}}_{\setminus i}^t = \frac{m \bar{\mathbf{w}}^t - \hat{\mathbf{w}}_i^t}{m-1}$.
 - 10: **end for**
 - 11: **Output:** $\bar{\mathbf{w}} = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$.
-

B. Reproducing Kernel Hilbert Space

When \mathcal{H} is a RKHS, that is $f(\mathbf{x}) = \sum_{j=1}^n w_j K(\mathbf{x}_j, \mathbf{x})$, we consider the following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^n} \frac{1}{n} \|\mathbf{K}_{\mathcal{S}_i} \mathbf{w} - \mathbf{y}_{\mathcal{S}_i}\|_2^2 + \lambda \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{w} + \frac{\gamma}{m-1} \sum_{j=1, j \neq i}^m \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j, \quad (12)$$

where $\mathbf{K}_{\mathcal{S}_i} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_{j'}}) \right]_{j, j'=1}^n$, $z_{t_j}, z_{t_{j'}} \in \mathcal{S}_i$, $\mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_k}) \right]_{j, k=1}^n$, $z_{t_j} \in \mathcal{S}_i, z_{t_k} \in \mathcal{S}_j$. Note that $\hat{\mathbf{w}}_i$ can be written as

$$\hat{\mathbf{w}}_i = \left(\underbrace{\frac{1}{n} \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n}_{:= \mathbf{A}_i} \right)^{-1} \left(\underbrace{\frac{1}{n} \mathbf{y}_{\mathcal{S}_i}}_{:= \mathbf{b}_i} - \frac{\gamma}{2} \bar{\mathbf{g}}_{\setminus i} \right),$$

where $\mathbf{g}_j = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j$ and $\bar{\mathbf{g}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{\mathbf{g}}_j$.

Similar with the linear space, we need to compute $\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i}$ in each iterative. From Lemma 3 (see in the last part), we know that

$$\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i} = \left(\bar{\mathbf{g}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i.$$

The Max-Discrepant Distributed Learning algorithm for RKHS is given in Algorithm 2. Compared with the traditional divide-and-conquer method, our MDD for RKHS only need add $\mathcal{O}(n)$ in each iteration for local machine.

Algorithm 2 Max-Discrepant Distributed Learning for RKHS (MDD-RKHS)

- 1: **Input:** λ, γ , kernel function K , \mathbf{X}, m, ζ .
 - 2: **For each worker node i :** $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$, and push $\hat{\mathbf{w}}_i^0$ to the server node.
 $// \mathbf{A}_i = \frac{1}{n} \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n, \mathbf{b}_i = \frac{1}{n} \mathbf{y}_{\mathcal{S}_i}$.
 - 3: **For server node:** $\hat{\mathbf{g}}_{i,j}^0 = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^0, i, j = 1, \dots, m, \bar{\mathbf{g}}_{\setminus i}^0 = \frac{m \bar{\mathbf{g}}^0 - \hat{\mathbf{g}}_i^0}{m-1}$.
 - 4: **for** $t = 1, 2, \dots$ **do**
 - 5: **For each worker node i :**
 Pull $\bar{\mathbf{g}}_{\setminus i}^{t-1}$ from server node.
 $\mathbf{d}_i^t = \left(\left(\bar{\mathbf{g}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) ./ \mathbf{b}_i, \hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$.
 - 6: Push $\hat{\mathbf{w}}_i^t$ to the server node.
 - 7: **For server node:**
 $\hat{\mathbf{g}}_{i,j}^t = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^t, i, j = 1, \dots, m, \bar{\mathbf{g}}_i^t = \frac{1}{m} \sum_{j=1}^m \hat{\mathbf{g}}_{i,j}^t$
 if $\frac{1}{m} \sum_{i=1}^m \|\bar{\mathbf{g}}_i^t - \bar{\mathbf{g}}_i^{t-1}\| \leq \zeta$ **end for**
 else $\bar{\mathbf{g}}_{\setminus i}^t = \frac{m \bar{\mathbf{g}}^t - \hat{\mathbf{g}}_i^t}{m-1}$.
 - 11: **end for**
 - 12: **Output:** $\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i$, where $\hat{f}_i = \mathbf{k}_{\mathcal{S}_i}^T \hat{\mathbf{w}}_i$, where $\mathbf{k}_{\mathcal{S}_i} = (K(\mathbf{x}_1, \cdot), \dots, K(\mathbf{x}_n, \cdot))^T, z_j \in \mathcal{S}_i$
-

Remark 2. From ensemble learning, to obtain good performance, the diversity of basis learning machines is very important [16]. The motivation of this paper was inspired by the ensemble learning, but one more thing should be

Table I

COMPARISON OF AVERAGE ROOT MEAN SQUARE ERROR OF OUR MDD-LS AND MDD-RKHS WITH RR, DRR, KRR, KRRR. WE BOLD THE NUMBERS OF THE BEST METHOD, AND UNDERLINE THE NUMBERS OF THE OTHER METHODS WHICH ARE NOT SIGNIFICANTLY WORSE THAN THE BEST ONE

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPred
RR	0.971	2.585	45.150	0.247	1.932	0.671	0.673	0.841	12.233
DRR-5	<u>0.989</u>	2.814	53.114	0.262	2.659	<u>0.681</u>	<u>0.680</u>	0.855	14.216
DRR-10	1.408	2.983	55.557	0.273	2.839	0.725	0.696	0.863	15.780
MDD-LS-5	<u>0.977</u>	<u>2.677</u>	46.184	<u>0.257</u>	<u>2.114</u>	<u>0.677</u>	<u>0.673</u>	<u>0.847</u>	<u>12.303</u>
MDD-LS-10	1.121	2.750	48.956	0.268	<u>2.352</u>	<u>0.703</u>	<u>0.685</u>	<u>0.854</u>	14.158
KRR	0.959	1.458	53.993	0.167	1.504	0.659	0.790	0.671	/
KDRR-5	<u>1.142</u>	2.389	<u>54.228</u>	0.419	1.598	0.873	0.866	<u>0.674</u>	<u>5.397</u>
KDRR-10	1.374	2.531	<u>56.233</u>	0.422	1.824	0.906	0.893	<u>0.687</u>	5.631
MDD-RKHS-5	<u>0.992</u>	2.030	<u>54.015</u>	0.214	<u>1.554</u>	<u>0.745</u>	<u>0.804</u>	<u>0.672</u>	5.350
MDD-RKHS-10	1.292	2.326	55.120	0.239	1.780	<u>0.773</u>	0.849	0.683	5.534

emphasized, the theoretical proof and algorithm design of this paper are not from the ensemble learning. The theoretical results w.r.t. diversity are given for a distributed setting has not seen before.

C. Complexity

Linear space: At the very beginning, we need $\mathcal{O}(nd^2)$ to compute the \mathbf{A}_i , $\mathcal{O}(d^3)$ to compute \mathbf{A}_i^{-1} for each worker node. In each iteration, worker nodes cost $\mathcal{O}(d)$ to compute \mathbf{d}_i^t and the server node costs $\mathcal{O}(md)$ to compute $\bar{\mathbf{w}}_i^t$. So, the sequential computation complexity is $\mathcal{O}(nd^2 + d^3 + Tmd)$, where T is the number of iteration. Moreover, the total communication complexity is $\mathcal{O}(Td)$.

RKHS: At the very beginning, we need $\mathcal{O}(n^2d)$ to compute the \mathbf{A}_i and $\mathcal{O}(n^3)$ to compute \mathbf{A}_i^{-1} . In each iteration, worker nodes cost $\mathcal{O}(n)$ to compute \mathbf{d}_i^t and the server node costs $\mathcal{O}(mn)$ to compute $\bar{\mathbf{g}}_i^t$. So, the sequential computation complexity is $\mathcal{O}(n^2d + n^3 + Tmn)$, where T is the number of iteration. Moreover, the total communication complexity is $\mathcal{O}(Tn)$.

Divide-and-conquer approach: The sequential complexities of linear space and RKHS are $\mathcal{O}(nd^2 + d^3)$ and $\mathcal{O}(n^2d + n^3)$, respectively. Meanwhile, the communication complexities are $\mathcal{O}(d)$ and $\mathcal{O}(n)$.

Global approach: The total complexities of linear space and RKHS are $\mathcal{O}(Nd^2 + d^3)$ and $\mathcal{O}(N^2d + N^3)$, respectively.

IV. EXPERIMENTS

In this part, we will compare our MDD methods with the global method and divide-and-conquer method in both Linear and RKHS Hypothesis. Actually, we compare six approaches: global Ridge Regression (RR) [17], divide-and-conquer Ridge Regression (DRR) and our MDD-LS (Algorithm 1) in Linear Hypothesis Space, meanwhile, global Kernel Ridge Regression (KRR) [18], divide-and-conquer Kernel Ridge Regression (KDRR) [2] and our MDD-RKHS (Algorithm 2) in Reproducing Kernel Hilbert Space. Based on the recent distributed machine learning platform PARAMETER SERVER [19], we implemented divide-and-conquer methods and MDD methods and do experiments on this framework.

We experiment on 10 publicly available datasets from LIBSVM data ¹. We run all methods on a computer node with 32 cores (2.40GHz) and 64 GB memory. While global methods only use a single CPU core, distributed methods use all cores to simulate parallel environment. For RKHS methods, we use the most popular Gaussian kernels

$$K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|_2^2}{2\sigma^2}\right)$$

as candidate kernels, and choose the best kernel from $\sigma \in \{2^i, i = -10, -9, \dots, 10\}$ by 5-folds cross-validation. The regularized parameterized $\lambda \in \{10^i, i = -6, -5, \dots, 3\}$ in all methods and $\gamma \in \{10^i, i = -6, -5, \dots, 3\}$ in MDD methods are determined by 5-folds cross-validation on training data. On each data set, we operate all methods 30 times with random partitions on all data sets of non-overlapping 70% training data and 30% testing data. All statements of statistical significance in the remainder refer to a 95% level of significance under t -test.

The root mean square error of all methods is reported in Table I. Meanwhile, we repeat distributed methods on different amount of worker nodes, 5 and 10 for simplification. Table I can be summarized as follows:

- 1) Our MDD-LS and MDD-RKHS exhibit better prediction accuracy than the DRR and KDRR over almost all data sets. This demonstrates the advantage of MDD methods in generalization performance.
- 2) Our MDD-LS and MDD-RKHS give comparable result with global methods on most of data sets.
- 3) Kernel methods can usually get more optimal results than linear methods do;
- 4) Some data sets are sensitive to data partition, whose results existing huge gap between global methods and distributed methods, such as space_ga and phishing for RKHS, while others are not.
- 5) The increase of worker nodes causes higher root mean square error.

The running time is reported in Table II, which can be summarized as follows:

- 1) Global methods cost more time than distributed methods do on all data sets.

¹Available at <https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/>

Table II
COMPARISON OF RUN TIME (SECOND) AMOUND OUR PROPOSED MDD-LS AND MDD-RKHS WITH OTHER METHODS.

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPred
RR	2.069	0.280	1.218	1.526	0.490	2.544	2.957	1.866	10.433
DRR-5	0.849	0.094	0.463	0.625	0.363	0.773	0.881	0.736	3.709
DRR-10	0.623	0.193	0.298	0.350	0.214	0.401	0.503	0.435	2.645
MDD-LS-5	0.875	0.115	0.587	0.664	0.427	1.208	1.167	0.876	5.474
MDD-LS-10	0.656	0.134	0.315	0.395	0.269	0.651	0.628	0.412	3.156
KRR	3.450	1.508	9.801	12.08	76.99	15.33	16.103	137.6	/
KDRR-5	1.487	0.295	3.374	1.451	5.524	6.021	5.913	40.22	86.754
KDRR-10	0.983	0.183	1.863	0.689	0.302	3.670	3.544	23.64	46.197
MDD-RKHS-5	1.692	0.331	5.637	1.901	29.85	8.628	9.454	73.09	167.208
MDD-RKHS-10	1.041	0.206	3.024	0.984	17.78	4.125	5.679	40.23	89.312

- 2) Kernel methods always spend more time than linear methods, because of higher computation complexity.
- 3) Distributed methods lead great speedup on some data sets.
- 4) The running time of distributed methods decays almost linearly associated with the increase of worker nodes.
- 5) Compared with global methods, our MDD methods own higher computational efficiency, while existing small distance away from divide-and-conquer methods.

The above results show that MDD methods need a bit more training time but make the performance gap between global methods and traditional distributed methods tighter, which is consistent with our theoretical analysis.

V. CONCLUSION

We studied the generalization performance of distributed learning, and derived a sharper generalization error bound, which is much sharper than existing generalization bounds of divide-and-conquer based distributed learning. Then, we designed two algorithms with statistical guarantees and fast convergence rates for linear space and RKHS: MDD-LS and MDD-RKHS. As we see from theoretical analysis and empirical results, our MDD is highly competitive with the existing divide-and-conquer methods, in terms of both practical performance and computational cost. Based on max-diversity of each local estimate, our analysis can be used as a solid basis for the design of new distributed learning algorithms.

VI. PROOF

A. The Key Idea

From the η -strongly convex of $R(f)$ of equation (5), we can obtain that

$$\begin{aligned} R(\bar{f}) &= R\left(\frac{1}{m} \sum_{i=1}^m \hat{f}_i\right) \\ &\leq \frac{1}{m} \sum_{i=1}^m R(\hat{f}_i) - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2. \end{aligned}$$

Therefore, we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{1}{m} \sum_{i=1}^m [R(\hat{f}_i) - R(f_*)] \\ &\quad - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2. \end{aligned} \quad (13)$$

In the next, we will estimate $R(\hat{f}_i) - R(f_*)$, which is built upon the following inequality from (4):

$$\begin{aligned} R(\hat{f}_i) - R(f_*) &+ \frac{\eta}{2} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2 \\ &\leq \langle \nabla R(\hat{f}_i), \hat{f}_i - f_* \rangle_{\mathcal{H}} \\ &= \langle \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}_i(\hat{f}_i) - \nabla \hat{R}_i(f_*)], \hat{f}_i - f_* \rangle_{\mathcal{H}} \\ &\quad + \langle \nabla R(f_*) - \nabla \hat{R}_i(f_*), \hat{f}_i - f_* \rangle_{\mathcal{H}} + \langle \nabla \hat{R}_i(\hat{f}_i), \hat{f}_i - f_* \rangle_{\mathcal{H}}. \end{aligned} \quad (14)$$

According to the convexity of $\hat{R}_i(\cdot)$ and the optimality condition of \hat{f}_i [20], we have

$$\langle \nabla \hat{R}_i(\hat{f}_i), \hat{f}_i - f_* \rangle_{\mathcal{H}} \geq 0, \forall f \in \mathcal{H}. \quad (15)$$

Substituting (15) into (14), we have

$$\begin{aligned} R(\hat{f}_i) - R(f_*) &+ \frac{\eta}{2} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2 \\ &\leq \langle \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}_i(\hat{f}_i) - \nabla \hat{R}_i(f_*)], \hat{f}_i - f_* \rangle_{\mathcal{H}} \\ &\quad + \langle \nabla R(f_*) - \nabla \hat{R}_i(f_*), \hat{f}_i - f_* \rangle_{\mathcal{H}} \\ &\leq \left(\underbrace{\left\| \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}_i(\hat{f}_i) - \nabla \hat{R}_i(f_*)] \right\|}_{:=A_1} \right. \\ &\quad \left. + \underbrace{\left\| \nabla R(f_*) - \nabla \hat{R}_i(f_*) \right\|}_{:=A_2} \right) \|\hat{f}_i - f_*\|. \end{aligned} \quad (16)$$

B. Proof of Theorem 1

To prove Theorem 1, we first give the following two lemmas (the proofs are given at the last part of this section).

Lemma 1. Under Assumptions 3 and 4, $\forall f \in \mathcal{N}(\mathcal{H}, \epsilon)$, at least $1 - \delta$, we have

$$\begin{aligned} &\left\| \nabla R(f) - \nabla R(f_*) - [\nabla \hat{R}_i(f) - \nabla \hat{R}_i(f_*)] \right\| \\ &\leq \frac{(\tau + \tau') \|f - f_*\| \log C(\mathcal{H}, \epsilon)}{n} \\ &\quad + \sqrt{\frac{(\tau + \tau')(R(f) - R(f_*)) \log C(\mathcal{H}, \epsilon)}{n}}. \end{aligned} \quad (18)$$

Lemma 2. Under Assumptions 3, at least $1 - \delta$, we have

$$\left\| \nabla R(f_*) - \nabla \hat{R}_i(f_*) \right\| \leq \frac{2M \log(2/\delta)}{n} + \sqrt{\frac{8\tau H_* \log(2/\delta)}{n}}, \quad (19)$$

where $H_* = \mathbb{E}_z [\ell(f_*, z)]$.

Proof of Theorem 1. According to the property of ϵ -net, one can see that there exists a point $\tilde{f} \in \mathcal{N}(\mathcal{H}, \epsilon)$ such that

$$\|\hat{f}_i - \tilde{f}\| \leq \epsilon.$$

According to **Assumptions 3** and **4**, we know that $R(f)$ and $\hat{R}(f)$ are both $(\tau + \tau')$ -smooth. Thus, we have

$$\begin{aligned} & \left\| \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}_i(\hat{f}_i) - \nabla \hat{R}_i(f_*)] \right\| \\ & \leq \left\| \nabla R(\tilde{f}) - \nabla R(f_*) - [\nabla \hat{R}_i(\tilde{f}) - \nabla \hat{R}_i(f_*)] \right\| + 2(\tau + \tau')\epsilon \\ & \stackrel{(18)}{\leq} \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \|\tilde{f} - f_*\|}{n} \\ & \quad + \sqrt{\frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\tilde{f}) - R(f_*))}{n}} + 2(\tau + \tau')\epsilon \\ & \leq \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \quad + \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \epsilon}{n} + 2(\tau + \tau')\epsilon \\ & \quad + \sqrt{\frac{\beta \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n}} \\ & \quad + \sqrt{\frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \left(|R(\hat{f}_i) - R(\tilde{f})| \right)}{n}} \\ & \stackrel{(8)}{\leq} \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \quad + \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \epsilon}{n} + 2(\tau + \tau')\epsilon \\ & \quad + \sqrt{\frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n}} \\ & \quad + \sqrt{\frac{(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n}} \end{aligned} \quad (20)$$

Substituting Eqs.(20) and (19) into Eq.(17), at least $1 - 2\delta$, we can obtain that

$$\begin{aligned} & R(\hat{f}_i) - R(f_*) + \frac{\eta}{2} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2 \\ & \leq \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \|\hat{f}_i - f_*\|_{\mathcal{H}}^2}{n} \\ & \quad + \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \epsilon \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \quad + 2(\tau + \tau')\epsilon \|\hat{f}_i - f_*\|_{\mathcal{H}} \\ & \quad + \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n}} \\ & \quad + \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n}} \\ & \quad + \frac{2M \log(2/\delta) \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \quad + \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{8\tau H_* \log(2/\delta)}{n}}. \end{aligned} \quad (21)$$

Note that

$$\sqrt{ab} \leq \frac{a}{2c} + \frac{bc}{2}, \forall a, b, c \geq 0.$$

Therefore, we can obtain that

$$\begin{aligned} & \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n}} \\ & \leq \frac{2(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n\eta} + \frac{\eta}{8} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2; \\ & \quad \frac{2M \log(2/\delta) \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \leq \frac{8M \log(2/\delta)}{n^2\eta} + \frac{\eta}{16} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2; \\ & \quad \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{8\eta H_* \log(2/\delta)}{n}} \\ & \leq \frac{64\eta H_* \log(2/\delta)}{n\eta} + \frac{\eta}{32} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2; \\ & \quad 2(\tau + \tau')\epsilon \|\hat{f}_i - f_*\|_{\mathcal{H}} \\ & \leq \frac{32(\tau + \tau')^2 \epsilon^2}{\eta} + \frac{\eta}{64} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2; \\ & \quad \|\hat{f}_i - f_*\|_{\mathcal{H}} \sqrt{\frac{(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n}} \\ & \leq \frac{32(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n\eta} + \frac{\eta}{128} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2; \\ & \quad \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \epsilon \|\hat{f}_i - f_*\|_{\mathcal{H}}}{n} \\ & \leq \frac{32(\tau + \tau') \log^2 C(\mathcal{H}, \epsilon) \epsilon^2}{n^2\eta} + \frac{\eta}{128} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2. \end{aligned}$$

Substituting the above inequation into (21), we can obtain that

$$\begin{aligned} & R(\hat{f}_i) - R(f_*) + \frac{\eta}{4} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2 \\ & \leq \frac{(\tau + \tau') \log C(\mathcal{H}, \epsilon) \|\hat{f}_i - f_*\|_{\mathcal{H}}^2}{n} \\ & \quad + \frac{2(\tau + \tau') \log C(\mathcal{H}, \epsilon) (R(\hat{f}_i) - R(f_*))}{n\eta} + \frac{8M \log(2/\delta)}{n^2\eta} \\ & \quad + \frac{64\tau H_* \log(2/\delta)}{n\eta} + \frac{32(\tau + \tau')^2 \epsilon^2}{\eta} \\ & \quad + \frac{32(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n\eta} \\ & \quad + \frac{32(\tau + \tau') \log^2 C(\mathcal{H}, \epsilon) \epsilon^2}{n^2\eta} \\ & \stackrel{(9)}{\leq} \frac{\eta}{4} \|\hat{f}_i - f_*\|_{\mathcal{H}}^2 + \frac{1}{2} (R(\hat{f}_i) - R(f_*)) + \frac{8(\tau + \tau') \log(2/\delta)}{n^2\eta} \\ & \quad + \frac{64\tau H_* \log(2/\delta)}{n\eta} + \frac{32(\tau + \tau')^2 \epsilon^2}{\eta} \\ & \quad + \frac{32(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n\eta} \\ & \quad + \frac{32(\tau + \tau') \log^2 C(\mathcal{H}, \epsilon) \epsilon^2}{n^2\eta}. \end{aligned}$$

Thus, with $1 - 2\delta$, we have

$$\begin{aligned} & R(\hat{f}_i) - R(f_*) \\ & \leq \frac{16M \log(2/\delta)}{n^2 \eta} + \frac{128\tau H_* \log(2/\delta)}{n\eta} \\ & \quad + \frac{32(\tau + \tau')^2 \epsilon^2}{\eta} + \frac{64(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n\eta} \\ & \quad + \frac{64(\tau + \tau') \log^2 C(\mathcal{H}, \epsilon) \epsilon^2}{n^2 \eta}. \end{aligned} \quad (22)$$

Combining (13) and (22), with $1 - \delta$, we have

$$\begin{aligned} & R(\bar{f}) - R(f_*) \\ & \leq \frac{16M \log(4m/\delta)}{n^2 \eta} + \frac{128\tau H_* \log(4m/\delta)}{n\eta} \\ & \quad + \frac{32(\tau + \tau')^2 \epsilon^2}{\eta} + \frac{64(\tau + \tau') L \log C(\mathcal{H}, \epsilon) \epsilon}{n\eta} \\ & \quad + \frac{64(\tau + \tau') \log^2 C(\mathcal{H}, \epsilon) \epsilon^2}{n^2 \eta} \\ & \quad - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2. \end{aligned}$$

□

C. Proof of Lemma 1

Proof. According to **Assumption 3** and 4, we know that $\nu(f, \cdot) = \nu(f, z) = \ell(f, z) + r(f)$ is $(\tau + \tau')$ -smooth, so we have

$$\|\nabla \nu(f, \cdot) - \nabla \nu(f_*, \cdot)\|_{\mathcal{H}} \leq (\tau + \tau') \|f - f_*\|_{\mathcal{H}} \quad (23)$$

Because $\nu(f, \cdot)$ is $(\tau + \tau')$ -smooth and convex, by (2.1.7) of [21], $\forall z \in \mathcal{Z}$, we have

$$\begin{aligned} & \|\nabla \nu(f, z) - \nabla \nu(f_*, z)\|^2 \\ & \leq (\tau + \tau') (\nu(f, z) - \nu(f_*, z) - \langle \nabla \nu(f_*, z), f - f_* \rangle_{\mathcal{H}}). \end{aligned}$$

Taking expectation over both sides, we can obtain that

$$\begin{aligned} & \mathbb{E}_{z \sim \mathbb{P}} [\|\nabla \nu(f, \cdot) - \nabla \nu(f_*, \cdot)\|^2] \\ & \leq (\tau + \tau') \left(R(\hat{f}_i) - R(f_*) - \langle \nabla R(f_*), f - f_* \rangle_{\mathcal{H}} \right) \\ & \leq (\tau + \tau') \left(R(\hat{f}_i) - R(f_*) \right) \end{aligned}$$

where the last inequality follows from the optimality condition of f_* , i.e.,

$$\langle \nabla R(f_*), f - f_* \rangle_{\mathcal{H}} \geq 0, \forall f \in \mathcal{H}.$$

Following Lemma 1 in [10], we have

$$\begin{aligned} & \left\| \nabla R(f) - \nabla R(f_*) - [\nabla \hat{R}_i(f) - \nabla \hat{R}_i(f_*)] \right\|_{\mathcal{H}} \\ & = \left\| \nabla R(f) - \nabla R(f_*) - \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} [\nabla \nu(f, z_i) - \nabla \nu(f_*, z_i)] \right\|_{\mathcal{H}} \\ & \leq \frac{2(\tau + \tau') \|f - f_*\|_{\mathcal{H}} \log(2/\delta)}{n} \\ & \quad + \sqrt{\frac{2(\tau + \tau') (R(f) - R(f_*)) \log(2/\delta)}{n}}. \end{aligned}$$

Taking the union bound over all $f \in \mathcal{N}(\mathcal{H}, \epsilon)$, we can prove Lemma 1. □

D. Appendix: Proof of Lemma 2

Proof. Since $\ell(f, \cdot)$ is η -smooth and nonnegative, from Lemma 4 of [22], we have

$$\|\nabla \ell(f_*, z_i)\|^2 \leq 4(\tau + \tau') \ell(f_*, z_i)$$

and thus

$$\begin{aligned} \mathbb{E}_{z \sim \mathbb{P}} [\|\nabla \ell(f_*, z)\|^2] & \leq 4(\tau + \tau') \mathbb{E}_{z \sim \mathbb{P}} [\ell(f_*, z)] \\ & = 4(\tau + \tau') R(f_*). \end{aligned}$$

From the **Assumption**, we have $\nabla \ell(f_*, z) \leq M, \forall z \in \mathcal{Z}$. Let $H(f) = R(f) - r(f)$ and $\hat{H}(f) = \hat{R}(f) - r(f)$. Then, according to Lemma 1 in [10], we know that

$$\begin{aligned} & \left\| \nabla R(f_*) - \nabla \hat{R}_i(f_*) \right\| = \left\| \nabla H(f_*) - \nabla \hat{H}_i(f_*) \right\| \\ & = \left\| \nabla H(f_*) - \frac{1}{n} \sum_{z_j \in \mathcal{S}_i} \nabla \ell(f_*, z_j) \right\| \\ & \leq \frac{2(\tau + \tau') \log(2/\delta)}{n} + \sqrt{\frac{8(\tau + \tau') H_* \log(2/\delta)}{n}}. \end{aligned}$$

□

E. Proof of Lemma 3

Lemma 3. For all $\ell \geq 1$, If $\mathbf{A} \in \mathbb{R}^{\ell \times \ell}$ is a symmetric matrix and $\mathbf{b}, \mathbf{d} \in \mathbb{R}^{\ell}$, $\mathbf{c} = \mathbf{A}^{-1} \mathbf{b} \in \mathbb{R}^{\ell}$, then we have

$$\mathbf{A}^{-1} \mathbf{d} = (\mathbf{d}^T \mathbf{c}) ./ \mathbf{b},$$

where $a ./ \mathbf{c} = (a/c_1, \dots, a/c_{\ell})^T$.

Proof. Since \mathbf{A} a symmetric matrix, we have

$$(\mathbf{A}^{-1} \mathbf{d})^T \mathbf{b} = \mathbf{d}^T \mathbf{A}^{-1} \mathbf{b} = \mathbf{d}^T \mathbf{c}.$$

Therefore, we can obtain that $\mathbf{A}^{-1} \mathbf{d} = (\mathbf{d}^T \mathbf{c}) ./ \mathbf{b}$. □

REFERENCES

- [1] Z.-H. Zhou, N. V. Chawla, Y. Jin, and G. J. Williams, "Big data opportunities and challenges: Discussions from data analytics perspectives [discussion forum]," *IEEE Computational Intelligence Magazine*, vol. 9, no. 4, pp. 62–74, 2014.
- [2] Y. Zhang, J. Duchi, and M. Wainwright, "Divide and conquer kernel ridge regression," in *Proceedings of Conference on Learning Theory (COLT 2013)*, 2013, pp. 592–617.
- [3] S.-B. Lin, X. Guo, and D.-X. Zhou, "Distributed learning with regularized least squares," *The Journal of Machine Learning Research*, vol. 18, no. 1, pp. 3202–3232, 2017.
- [4] Y. Zhang, M. J. Wainwright, and J. C. Duchi, "Communication-efficient algorithms for statistical optimization," in *Advances in Neural Information Processing Systems*, 2012, pp. 1502–1510.
- [5] D. Gillick, A. Faria, and J. DeNero, "Mapreduce: Distributed computing for machine learning," *Berkley, Dec*, vol. 18, 2006.
- [6] X. Wu, X. Zhu, G.-Q. Wu, and W. Ding, "Data mining with big data," *IEEE transactions on knowledge and data engineering*, vol. 26, no. 1, pp. 97–107, 2014.
- [7] E. D. Vito, A. Caponnetto, and L. Rosasco, "Model selection for regularized least-squares algorithm in learning theory," *Foundations of Computational Mathematics*, vol. 5, no. 1, pp. 59–85, 2005.
- [8] A. Caponnetto and E. D. Vito, "Optimal rates for the regularized least-squares algorithm," *Foundations of Computational Mathematics*, vol. 7, no. 3, pp. 331–368, 2007.
- [9] I. Steinwart, D. Hus, and C. Scovel, "Optimal rates for regularized least squares regression," in *Proceedings of the Conference on Learning Theory (COLT 2009)*, 2009.

- [10] S. Smale and D.-X. Zhou, "Learning theory estimates via integral operators and their approximations," *Constructive approximation*, vol. 26, no. 2, pp. 153–172, 2007.
- [11] I. Steinwart and A. Christmann, *Support vector machines*. New York: Springer Verlag, 2008.
- [12] Z.-C. Guo, S.-B. Lin, and D.-X. Zhou, "Learning theory of distributed spectral algorithms," *Inverse Problems*, vol. 33, no. 7, p. 074009, 2017.
- [13] L. Zhang, T. Yang, and R. Jin, "Empirical risk minimization for stochastic convex optimization: $O(1/n)$ -and $O(1/n^2)$ -type of risk bounds," in *Proceedings of the Conference on Learning Theory (COLT 2017)*, 2017, pp. 1954–1979.
- [14] G. Pisier, *The volume of convex bodies and Banach space geometry*. Cambridge University Press, 1999, vol. 94.
- [15] D.-X. Zhou, "The covering number in learning theory," *Journal of Complexity*, vol. 18, no. 3, pp. 739–767, 2002.
- [16] Z.-H. Zhou, *Ensemble methods: foundations and algorithms*. Chapman and Hall/CRC, 2012.
- [17] A. E. Hoerl and R. W. Kennard, "Ridge regression: Biased estimation for nonorthogonal problems," *Technometrics*, vol. 12, no. 1, pp. 55–67, 1970.
- [18] S. An, W. Liu, and S. Venkatesh, "Fast cross-validation algorithms for least squares support vector machine and kernel ridge regression," *Pattern Recognition*, vol. 40, no. 8, pp. 2154–2162, 2007.
- [19] M. Li, D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J. Shekita, and B.-Y. Su, "Scaling distributed machine learning with the parameter server," in *OSDI*, vol. 14, 2014, pp. 583–598.
- [20] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge university press, 2004.
- [21] Y. Nesterov, *Introductory Lectures on Convex Optimization: A Basic Course*. Kluwer Academic Publishers, 2004.
- [22] N. Srebro, K. Sridharan, and A. Tewari, "Optimistic rates for learning with a smooth loss," *arXiv preprint arXiv:1009.3896*, 2010.