

Automatic Classifiers as Scientific Instruments: One Step Further Away from Ground-Truth

Jacob Whitehill (jrwhitehill@wpi.edu)
Worcester Polytechnic Institute

Abstract—Automatic detectors of facial expression, gesture, affect, etc., can serve as scientific instruments to measure many behavioral and social phenomena (e.g., emotion, empathy, stress, engagement, etc.), and this has great potential to advance basic science. However, when a detector d is trained to approximate an existing measurement tool (e.g., observation protocol, questionnaire), then care must be taken when interpreting measurements collected using d since they are one step further removed from the underlying construct. We examine how the accuracy of d , as quantified by the correlation q of d 's outputs with the ground-truth construct U , impacts the estimated correlation between U (e.g., stress) and some other phenomenon V (e.g., academic performance). In particular: (1) We show that if the true correlation between U and V is r , then the expected sample correlation, over all vectors \mathcal{T}^n whose correlation with U is q , is qr . (2) We derive a formula to compute the probability that the sample correlation (over n subjects) using d is positive, given that the true correlation between U and V is negative (and vice-versa). We show that this probability is non-negligible (around 10 – 15%) for values of n and q that have been used in recent affective computing studies. (3) With the goal to reduce the variance of correlations estimated by an automatic detector, we show empirically that training multiple neural networks $d^{(1)}, \dots, d^{(m)}$ using different training configurations (e.g., architectures, hyperparameters) for the same detection task provides only limited “coverage” of \mathcal{T}^n .

I. INTRODUCTION

Automatic classifiers have the potential to advance basic research in psychology, education, medicine, and many other fields by serving as *scientific instruments* that can measure behavioral, medical, social, and other phenomena with higher temporal resolution, lower cost, and greater consistency than is possible with traditional methods such as human-coded questionnaires or observation protocols. The affective computing community is starting to see some first fruits of this potential: Perugia, et al. [15] used the Empatica E4 wristband sensor to explore the relationship between participants' ($n = 14$) electrodermal activity (EDA) and their emotional states when playing cognitive games. Parra, et al. [14] used the Emotient facial expression recognition software to identify a positive correlation ($r = 0.32$, $n = 59$ participants) between emotions and adult attachment [3]. Chen, et al. [2] used Emotient in a study of how facial emotion is associated with job interview performance among $n = 4$ participants.

In most empirical studies designed to assess the possible relationship between two phenomena U and V (e.g., engagement [11], grit [4], stress, attachment [3], academic performance, etc.), the investigator chooses a validated instrument for each phenomenon and records measurements

of each variable for n participants. She/he then computes a statistic, such as the Pearson product-moment coefficient, that captures the strength and polarity (as well as statistical significance) of the relationship between the two variables. Machine learning, and automatic face and gesture recognition in particular, offers the potential to create a new array of scientific instruments with important advantages compared to standard measurement tools. However, they also bring a potential pitfall that – while not fundamentally new, i.e., there is always a separation between a construct and its measurement – is exacerbated compared to using standard measurements: If one creates a new scientific instrument by training an automatic detector d to mimic a standard instrument as closely as possible, then the scientific instrument d is *one degree of separation further removed from the underlying phenomenon U* – i.e., it is an estimator of another estimator.

Motivating example: Suppose a behavioral scientist wishes to examine the relationship between stress (construct U) and academic performance (construct V). Using a traditional approach, she/he could conduct an experiment in which each participant completes some cognitively demanding task and then takes a test. To measure stress, the scientist could also ask each participant to complete an established survey, e.g., the Dundee State Stress Questionnaire [10]. The relationship between U and V could then be estimated as the correlation $r = \rho(\mathbf{u}, \mathbf{v})$ between the vector of test scores \mathbf{v} and the corresponding vector of stress measurements \mathbf{u} over all n participants.

However, suppose that the researcher also has access to an *automatic stress detector* d that uses the participant's face pixels to measure his/her stress level. Suppose that the accuracy of d was previously validated w.r.t. a standard stress questionnaire (like [10]), and the validation showed that the outputs of d , which we denote with $\hat{\mathbf{u}}$, have an expected correlation of q with the standard questionnaire. What could go wrong, in terms of spurious deductions, when the correlation r between U and V is estimated as $\rho(\hat{\mathbf{u}}, \mathbf{v})$ instead of $\rho(\mathbf{u}, \mathbf{v})$?

Figure 1 shows one hypothetical example of what can go wrong: vectors \mathbf{u} and \mathbf{v} contain measurements from n participants of constructs U and V , respectively, where \mathbf{u} is obtained through a standard instrument. The Pearson product-moment correlation between two vectors can be written as:

$$\rho(\mathbf{u}, \mathbf{v}) = \frac{(\mathbf{u} - \mu_{\mathbf{u}})^{\top} (\mathbf{v} - \mu_{\mathbf{v}})}{\|\mathbf{u} - \mu_{\mathbf{u}}\|_2 \|\mathbf{v} - \mu_{\mathbf{v}}\|_2}$$

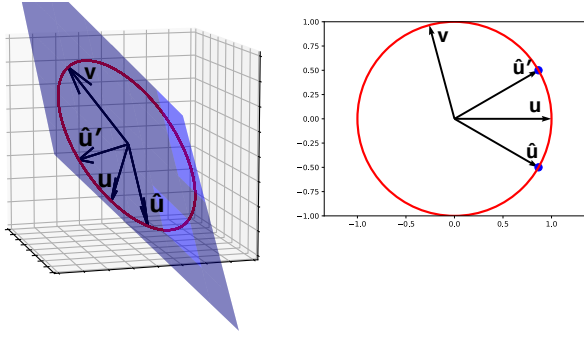


Fig. 1. \mathbf{u} and \mathbf{v} are measurements of some behavioral, social, educational, (or other) phenomena for n participants in an experiment. $\hat{\mathbf{u}}$ (or $\hat{\mathbf{u}}'$) are proxy measurements of \mathbf{u} that were obtained from an automatic detector. Both $\hat{\mathbf{u}}$ and $\hat{\mathbf{u}}'$ have the same correlation ($r \approx 0.867$) with \mathbf{u} . However, depending on *which* vector is obtained from the detector, the estimated correlation can be very different: $\rho(\hat{\mathbf{u}}, \mathbf{v}) < 0$, but $\rho(\hat{\mathbf{u}}', \mathbf{v}) > 0$.

where $\mu_{\mathbf{u}}$ (or $\mu_{\mathbf{v}}$) is a vector whose elements equal the mean value of \mathbf{u} (or \mathbf{v}). If \mathbf{u} and \mathbf{v} both have 0-mean and unit-length, then their correlation depends only on the *angle* between them:

$$\rho(\mathbf{u}, \mathbf{v}) = \mathbf{u}^\top \mathbf{v} = \cos \angle(\mathbf{u}, \mathbf{v})$$

In the figure, this correlation is $\cos(105^\circ) \approx -0.259$, i.e., the data suggest that U is *negatively* correlated with V . Suppose instead that the researcher had used an automatic detector d to obtain $\hat{\mathbf{u}}$, where previous analysis had established that the expected correlation of d 's outputs and the standard instrument was $q = \cos(30^\circ) \approx 0.866$. If the researcher uses the correlation $\rho(\hat{\mathbf{u}}, \mathbf{v})$ to estimate the relationship r between U and V , then she/he would obtain $\cos(135^\circ) = -0.707$ – a much larger magnitude, but at least the same sign as, the $r = -0.259$ obtained using a standard instrument for U . But the bigger problem is the following: $\hat{\mathbf{u}}$ is not the only vector whose correlation with the “ground-truth” measurements \mathbf{u} is q . Vector $\hat{\mathbf{u}}'$ also has the same correlation. If the researcher obtained measurements $\hat{\mathbf{u}}'$, then she/he would deduce a *positive* correlation of $\rho(\hat{\mathbf{u}}', \mathbf{v}) = \cos(75^\circ) \approx 0.259$ – this is opposite to the correlation obtained with a standard instrument.

In this paper we explore how the accuracy q of a scientific instrument d , as measured by the Pearson correlation with the ground-truth construct U , impacts the estimated correlation between constructs U and V . Although there are various ways of quantifying the relationship between two vectors of measurements (e.g., RMSE, MAE), the Pearson correlation is one of the most commonly used metrics. **Contributions:** (1) We prove that $\mathbb{E}[\rho(\hat{\mathbf{u}}, \mathbf{v})] = qr$, where $\hat{\mathbf{u}}$ is sampled uniformly over the $(n-3)$ -sphere \mathcal{T}^n of 0-mean unit-vectors whose correlation with \mathbf{u} is q . Next, as one of the most fundamental aspects of the relationship between two variables is whether they are positively or negatively correlated, (2) we derive

a function h to compute the probability that the sample correlation (over n subjects) using d is positive, given that the true correlation between U and V is negative (and vice-versa). We also prove that h is monotonically decreasing in n and in q , i.e., the danger of a false correlation is mitigated by training a more accurate detector or collecting data from more participants. Finally, (4) we explore to what extent the sphere \mathcal{T}^n can be “covered” by measurement vectors $\hat{\mathbf{u}}^{(1)}, \dots, \hat{\mathbf{u}}^{(m)}$ obtained by training m different neural networks trained on the same dataset for the same detection task but using different configurations (e.g., CNN versus MLP, hyperparameters, etc.).

II. RELATED WORK

The issue of how product-moment (Pearson) correlations among a subset of variables constrain the possible correlations among the remaining variables has interested statisticians since the 1960s. While there has been significant prior work on the trivariate case in particular, we are not aware of any work that proves exactly the same results as what we present here. Priest [17] showed a lower bound on the mean intercorrelation between variables. Glass and Collins [6], and also Leung and Lam [9], proved that, in trivariate distributions, there are range restrictions on the possible correlations between U and V when the correlations between V and W and between U and W are already known. Olkin [13] extended this result to multivariate distributions beyond 3 variables.

More recent, and most similar to our work, is a study by Carlson and Herdman [1] from the operations research community in 2012. They examined the methodological risk of using proxy measures to estimate the correlations between different constructs. Using analytical results by Leung and Lam [9], they show how the observed correlations between \mathbf{u} and $\hat{\mathbf{u}}$ can vary substantially as a function of the reliability of a proxy measure $\hat{\mathbf{u}}$ of \mathbf{u} . In contrast to our work, theirs is based on simulations and contains no formal proofs.

III. EXPECTED CORRELATION OF $\hat{\mathbf{u}} \in \mathcal{T}^n$ WITH \mathbf{v}

When we use an automatic face or gesture classifier d to obtain a vector of measurements, then we obtain a vector $\hat{\mathbf{u}}$ whose correlation with the underlying construct U (e.g., stress) is q . However, as illustrated in the example above, there can be multiple such vectors, and which one is obtained can make a big difference on the estimated correlation. As we show below, the set of 0-mean unit-length vectors with a fixed correlation to another unit-vector is an $(n-3)$ -sphere embedded in \mathbb{R}^n . If we sampled uniformly at random from this sphere, then what would be the expected sample correlation between $\hat{\mathbf{u}}$ and some other vector \mathbf{v} (e.g., academic performance)? To simplify our analyses below, we assume $\mathbf{u}, \hat{\mathbf{u}}, \mathbf{v}$ all have 0-mean and unit-length since Pearson correlation is invariant to these quantities. (Regarding the uniformity assumption: see Future Work in Section VI.)

Proposition 1: Let \mathbf{u}, \mathbf{v} be n -dimensional, 0-mean, unit-length vectors with a Pearson product-moment correlation $\rho(\mathbf{u}, \mathbf{v}) = r$. Then (1) the set \mathcal{T}^n of 0-mean, unit-length

vectors whose correlation with \mathbf{u} is q is an $(n-3)$ -sphere embedded in \mathbb{R}^n . Moreover, (2) if $\hat{\mathbf{u}}$ (typeset in Futura to denote a random variable) is a random vector sampled uniformly from \mathcal{T}^n , then the expected sample correlation $E[\rho(\hat{\mathbf{u}}, \mathbf{v})] = qr$.

Proof: The set of all 0-mean n -vectors constitutes a hyperplane

$$\mathcal{H} \doteq \{\mathbf{x} \in \mathbb{R}^n : \mathbf{1}^\top \mathbf{x} = 0\}$$

that passes through the origin with normal vector $\mathbf{1} \doteq (1, \dots, 1)$. The set of all unit-length vectors constitutes an $(n-1)$ -sphere

$$\mathcal{S}^{n-1} \doteq \{\mathbf{x} \in \mathbb{R}^n : \|\mathbf{x}\|_2 = 1\}$$

embedded in \mathbb{R}^n . Therefore, $\mathbf{u}, \mathbf{v}, \hat{\mathbf{u}} \in \mathcal{H} \cap \mathcal{S}^{n-1}$. Figure 1 (left) shows \mathcal{H} in blue, as well as the intersection of \mathcal{H} with \mathcal{S}^{n-1} as a red circle. Since all three vectors have 0-mean and unit-length, then the correlations between these vectors depend only on the angles between them. Hence, w.l.o.g. we can rotate the axes so that \mathcal{H} consists of all vectors whose first coordinate is 0, and all correlations will be preserved. After doing so, the only remaining constraint is that the projected vectors have unit-length.

More precisely, we can compute an orthonormal basis \mathbf{B} such that the first coordinate of vector $\mathbf{B}\mathbf{x}$ is 0 for every $\mathbf{x} \in \mathcal{H}$, and such that

$$\mathbf{B}\mathbf{u} = (0, 1, \overbrace{0, 0, \dots, 0}^{n-2}) \quad (1)$$

$$\mathbf{B}\mathbf{v} = (0, a, b, 0, \dots, 0) \quad (2)$$

Geometrically, this means that we can define \mathbf{B} so that the projected \mathbf{u}, \mathbf{v} lie in the plane spanned by the second and third vectors in basis \mathbf{B} (see Figure 1 (right)) – this makes the rest of the derivation much simpler. a represents the component of \mathbf{v} parallel to \mathbf{u} , and b is the component orthogonal to \mathbf{u} . Since the correlation between \mathbf{u} and \mathbf{v} is r , and since \mathbf{B} is orthonormal, then

$$\begin{aligned} (\mathbf{B}\mathbf{u})^\top (\mathbf{B}\mathbf{v}) &= \mathbf{u}^\top \mathbf{v} \\ &= r \\ &= 0 \times 0 + 1 \times a + 0 \times b + 0 + \dots + 0 \\ &= a \end{aligned}$$

and hence $a = r$. Since $\|\mathbf{v}\|_2 = \|\mathbf{B}\mathbf{v}\|_2 = 1$, then $b = \sqrt{1-r^2}$.

Now consider any vector $\hat{\mathbf{u}}$ whose correlation with \mathbf{u} is q . Let us define $(\hat{u}_1, \dots, \hat{u}_n) \doteq \mathbf{B}\hat{\mathbf{u}}$. By construction of \mathbf{B} , we already know that $\hat{u}_1 = 0$. We also have

$$\begin{aligned} (\mathbf{B}\hat{\mathbf{u}})^\top (\mathbf{B}\mathbf{u}) &= \hat{\mathbf{u}}^\top \mathbf{u} \\ &= q \\ &= \hat{u}_1 \times 0 + \hat{u}_2 \times 1 + \hat{u}_3 \times 0 + \dots + \hat{u}_n \times 0 \\ &= \hat{u}_2 \end{aligned}$$

and hence $\hat{u}_2 = q$. Since $\mathbf{B}\hat{\mathbf{u}}$ is a unit-vector, then

$$\mathbf{B}\hat{\mathbf{u}} \in \left\{ (0, q, \hat{u}_3, \dots, \hat{u}_n) : \sum_{i=3}^n \hat{u}_i^2 = 1 - q^2 \right\}$$

This set is the surface of an $(n-3)$ -sphere, with radius $\sqrt{1-q^2}$, embedded in \mathbb{R}^n . Since \mathbf{B} simply rotates the axes, then \mathcal{T}^n is likewise a $(n-3)$ -sphere embedded in \mathbb{R}^n . This proves part 1.

For part 2: When sampling uniformly from \mathcal{T}^n , the distribution of \hat{u}_3 on the $(n-3)$ -sphere is symmetrical about 0. Then $E[\hat{u}_3] = 0$, and hence:

$$\begin{aligned} E[\rho(\hat{\mathbf{u}}, \mathbf{u})] &= E[\hat{\mathbf{u}}^\top \mathbf{v}] \\ &= E[(\mathbf{B}\hat{\mathbf{u}})^\top (\mathbf{B}\mathbf{v})] \\ &= E[0 + q \times r + \hat{u}_3 \times \sqrt{1-r^2}] \\ &= qr + E[\hat{u}_3] \sqrt{1-r^2} \\ &= qr \end{aligned}$$

Example: For the case $n = 3$, consider the four vectors shown in Figure 1 (left) whose values are approximately:

$$\begin{aligned} \mathbf{u} &= (.816, -.408, -.408) & \mathbf{v} &= (-.211, -.577, .788) \\ \hat{\mathbf{u}} &= (.707, 0, -.707) & \hat{\mathbf{u}}' &= (.707, -.707, 0) \end{aligned}$$

By construction, $\rho(\hat{\mathbf{u}}, \mathbf{u}) = \rho(\hat{\mathbf{u}}', \mathbf{u}) = \cos(30^\circ) = q$, and $\rho(\mathbf{u}, \mathbf{v}) = \cos(105^\circ) = r$. Via a change of basis \mathbf{B} , the vectors can be rotated so that

$$\begin{aligned} \mathbf{B}\mathbf{u} &= (0, 1, 0) & \mathbf{B}\mathbf{v} &= (0, \cos(105^\circ), \sin(105^\circ)) \\ \mathbf{B}\hat{\mathbf{u}} &= (0, \cos(30^\circ), -\sin(30^\circ)) & \mathbf{B}\hat{\mathbf{u}}' &= (0, \cos(30^\circ), \sin(30^\circ)) \end{aligned}$$

The rotated vectors are shown in Figure 1 (right). The set \mathcal{T}^3 contains exactly two elements (since it is a 0-sphere): $\hat{\mathbf{u}}$ and $\hat{\mathbf{u}}'$. If $\hat{\mathbf{u}}$ is sampled uniformly at random from \mathcal{T}^3 , then $E[\rho(\hat{\mathbf{u}}, \mathbf{v})] = qr \approx -.224$. This result agrees with

$$\frac{1}{2} [\rho(\hat{\mathbf{u}}, \mathbf{u}) + \rho(\hat{\mathbf{u}}', \mathbf{u})] = \frac{1}{2} [\cos(135^\circ) + \cos(75^\circ)] \approx -.224$$

IV. PROBABILITY OF FALSE CORRELATIONS

One of the most fundamental distinctions is whether two phenomena are positively or negatively correlated with each other (or neither). What is the probability that $\rho(\hat{\mathbf{u}}, \mathbf{v}) \geq 0$ given that the correlation $r < 0$ (*false positive correlation*); or that $\rho(\hat{\mathbf{u}}, \mathbf{v}) < 0$ given that the correlation $r \geq 0$ (*false negative correlation*)? How do these probabilities change as n increases or q increases?

Proposition 2: Let $q \in (0, 1]$ be the correlation between the detector's output $\hat{\mathbf{u}}$ and ground-truth \mathbf{u} ; let r be the correlation between \mathbf{u} and \mathbf{v} ; and let $\hat{\mathbf{u}}$ be sampled uniformly from \mathcal{T}^n . If $r < 0$, then the probability of a *false positive* correlation (in the sense defined above) is given by the function

$$h(n, q, r) = \begin{cases} \frac{1}{2} \Pr[c^2 \leq 1 - q^2] & n = 3 \\ \frac{1}{2} \int_0^\infty f_1(t) F_{n-3} \left(\frac{1-q^2-c^2}{c^2} t \right) dt & n > 3 \end{cases}$$

where f_k and F_k are the PDF and CDF of a χ^2 -random variable with k degrees of freedom, respectively, and $c = |qr|/\sqrt{1-r^2}$. Similarly, if $r > 0$, then the probability of a *false negative* correlation is also given by h .

Proof: See appendix. ■

Proposition 3: For every fixed $c > 0$ and $q \in (0, 1]$, function h is monotonically decreasing in n .

Proof: See appendix. ■

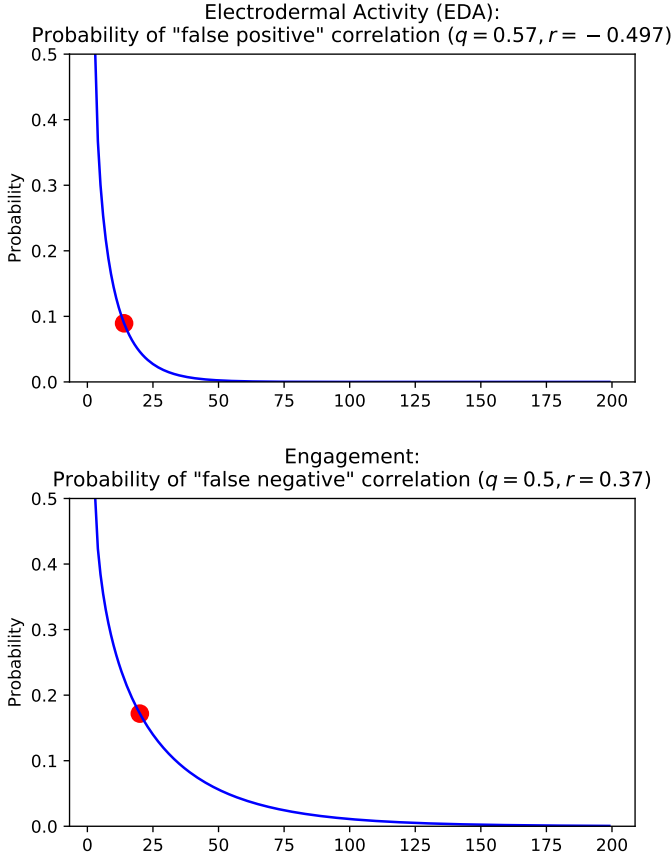


Fig. 2. Probability of a false correlation, for a fixed q and r , as a function of n . The probability decreases monotonically, but for small n it can still be substantial. The red dots indicate the n from two recent behavioral studies that used an automatic detector of affective state as a scientific instrument. **Top:** Example inspired by a study on electrodermal activity in [15]. **Bottom:** Example inspired by a study on student engagement in [19].

Proposition 4: For every fixed $n > 3$, function h is monotonically decreasing in $q \in (0, 1]$.

Proof: See appendix. ■

A. Case Studies

To put these theoretical results into perspective, we conducted simulations based on two recent affective computing studies that used automated detectors as scientific instruments. The first study ($n = 14$), by Perugia, et al. [15], used an Empatica E4 wrist sensor to investigate how electrodermal activity (EDA) (U) is correlated with the subjects' emotions (V). The second study ($n = 20$), by Whitehill, et al. [19], explored the relationship between student engagement (U), as measured by an engagement detector that analyzes static images of students' faces, and test performance (V) in a cognitive skills training task. In order to estimate the probability of a false correlation (in the sense described above), we need to know the accuracy of the automatic detector – i.e., the correlation q between the automatic measurements and ground-truth of construct U – as well as the *true* correlation r between constructs U and V .

Estimating q and r : The value of q can easily be estimated using cross-validation or other standard procedures. For the

first study (EDA), we use the value $q = 0.57$ reported in [16] for cognitive tasks with a distill forearm sensor of EDA. For r , we hypothesize that the ground-truth correlation between U and V is exactly what was estimated by the authors [15] (using the E4 sensor and emotion survey instruments): $r = -0.497$. For the second study (Engagement), we use the value $q = 0.50$ (subject-independent cross-validation correlation reported in [19]). For r , we use the correlation obtained by the authors ($r = 0.37$) when correlating test performance with *human*-labeled student engagement.

Results: Plots of the probability of a false correlation (obtained from function h derived above) as a function of the number of participants n are shown for each study (with their associated q and r values) in Figure 2. The red dot in each graph shows the actual number of participants from each experiment. While the probability is almost nil for $n \geq 200$, for more modest numbers of participants it is well above 0. For the values $n = 14$ and $n = 20$, these probabilities are non-negligible (around 10 – 15%). **The possibility of a false correlation is not protected against by statistical significance testing** – it is possible for the estimated correlation between constructs U and V to be highly significant and yet have the wrong sign compared to the ground-truth correlation. While this is almost always theoretically possible due to the inherent separation between a construct and its measurement, the use in basic research of automatic detectors that are trained to estimate another estimator can make this problem worse.

V. COVERAGE OF \mathcal{T}^n WHEN TRAINING A DETECTOR

Given that *which* vector $\hat{u} \in \mathcal{T}^n$ is obtained from an automatic detector d can substantially impact the estimated correlation $\rho(\hat{u}, \mathbf{v})$ between constructs U and V , it could be useful to *average* the sample correlations over *many* vectors $\hat{u}^{(1)}, \dots, \hat{u}^{(m)}$ from \mathcal{T}^n , i.e., to compute $\frac{1}{m} \sum_{i=1}^m \rho(\hat{u}^{(i)}, \mathbf{v})$. This could help to reduce the variance of the estimator. In this section we explore whether it is feasible to generate many different $\hat{u}^{(1)}, \dots, \hat{u}^{(m)}$, all with similar correlation q with ground-truth \mathbf{u} , by training a set of automatic detectors $d^{(1)}, \dots, d^{(m)}$ using slightly different training configurations. In particular, we varied: (1) the architecture (convolutional versus feed-forward MLP) and (2) the random seed of weight initialization. Inspired by recent work by Huang, et al. [7] on how an entire ensemble of detectors can be created during a *single* training run, we also varied (3) the number of training epochs, and saved snapshots of the trained detectors at regular intervals.

During training, each detector's estimates \hat{u} of the test labels evolves, and so does the correlation between \hat{u} and the ground-truth labels \mathbf{u} . However, for the tasks we examined (described below), the test correlations tend to stabilize over time, and they converge to roughly the same value even across different training runs and detection architectures. Suppose the average correlation (across all training configurations) of the machine's outputs \hat{u} with the test labels \mathbf{u} is q . Then we can collect all the $\hat{u}^{(1)}, \dots, \hat{u}^{(m)}$ (produced by




Fig. 3. Examples of face images used in the experiments in Section V-E. **Top:** student engagement dataset [19], in which engagement is rated on a 1-4 scale. **Bottom:** images from the GENKI [8] dataset that were labeled with their perceived age (in years).

detectors $d^{(1)}, \dots, d^{(m)}$ and project them onto the $(n-3)$ -sphere \mathcal{T}^n of 0-mean, unit-length vectors whose correlation with \mathbf{u} is q . We can then visualize the “coverage” of this sphere by projecting it onto a 2-D plane, and we can compare the coverage to a random sample of m elements of \mathcal{T}^n .

We explored the coverage of \mathcal{T}^n for two automatic face analysis problems: student engagement recognition and age estimation. For training and testing, we used the HBCU [19] (Engagement) and GENKI [8] (Age) datasets, respectively; see Figure 3 for labeled examples.

A. Detection architectures

For both applications, we analyzed cropped grayscale 48×48 face images and regressed a positive number. We implemented a CNN consisting of:

Input – Conv(16, 5) – ReLU – MP – BN – Conv(32, 5)
– ReLU – MP – BN – FC(16) – BN – ReLU – FC(1)

where Conv(c, k) is a convolutional layer with c output channels and a $k \times k$ spatial kernel of stride 1, and MP is a max-pooling layer of stride 2.

For engagement, we also implemented an MLP:

Input – FC(100) – ReLU – BN – FC(1) – ReLU

where FC(m) means fully-connected layer with m neurons, and BN is batch normalization. (We found during pilot experimentation that this approach did not generalize well for age estimation).

We note that our focus was not on obtaining state-of-the-art accuracy for these detection tasks. Rather, we wished to explore, for a plausible neural network-based detector, how much variation in vectors $\hat{\mathbf{u}}$ we can obtain by varying the hyperparameters.

B. Training procedures

Engagement detector: We performed 25 training runs for each of the two network architectures (CNN, MLP). Training data consisted of 7629 face images from 15 subjects of HBCU [19], and testing data were 500 images from the remaining 5 subjects. (This corresponds to just one cross-validation fold from the original study [19].) Optimization was performed using SGD for 10000 iterations, and the

network weights were saved every 1000 iterations. In total, this produced 500 detectors. The average correlation (over all 500 detectors) between the detectors’ test outputs $\hat{\mathbf{u}}$ and ground-truth \mathbf{u} was 0.55 (s.d. 0.079). Inspired by [7], we also tried both cosine and triangular [18] learning rates. However, in pilot testing we found that these delivered worse accuracy than exponential learning rate decay and we abandoned the approach.

Age detector: We performed 25 training runs for the CNN using SGD for 10000 with snapshots every 1000 iterations, as for engagement recognition. This produced 250 detectors. Training data consisted of 31040 face images of the GENKI dataset [8], and testing data consisted of 500 face images. The average correlation of the automatic measurements with ground-truth on the test set was 0.55 (s.d. 0.11).

C. Visualizing elements of the $(n-3)$ -sphere \mathcal{T}^n

Given a set of trained detectors and corresponding age/engagement estimates $\hat{\mathbf{u}}^{(1)}, \dots, \hat{\mathbf{u}}^{(m)} \in \mathbb{R}^n$ whose correlation with ground-truth \mathbf{u} is approximately q , we can visualize how these vectors “cover” the $(n-3)$ -sphere \mathcal{T}^n using the following procedure:

- 1) Normalize \mathbf{u} , as well as each $\hat{\mathbf{u}}^{(j)}$, to have 0-mean and unit-length.
- 2) Compute an orthonormal basis \mathbf{B} (e.g., using a QR decomposition) so that (a) the first component of $\mathbf{B}\mathbf{x}$ is 0 for every $\mathbf{x} \in \mathcal{H}$, and (b) $\mathbf{B}\mathbf{u} = (0, 1, 0, 0, \dots, 0)$ (see Equation 1).
- 3) Project each $\hat{\mathbf{u}}^{(j)}$ onto the new basis \mathbf{B} . By construction, the first component of each projection will be 0 and the second component will be $q = \rho(\hat{\mathbf{u}}^{(j)}, \mathbf{u})$.
- 4) Define each $\mathbf{x}^{(j)}$ to be the last $n-2$ components of vector $\mathbf{B}\hat{\mathbf{u}}^{(j)}$.
- 5) Project the $\{\mathbf{x}^{(j)}\}$ onto the two principal axes obtained from principal component analysis (PCA).

Since the 2-D projection of a 0-centered sphere onto any orthonormal projection is a disc, the output of the procedure above is a set of points that lie on a disc of radius $\sqrt{1-q^2}$.

D. Generating random vectors on \mathcal{T}^n

In order to assess how evenly the sphere \mathcal{T}^n is “covered” by the vectors obtained from the automatic detectors, we can generate random vectors of \mathcal{T}^n and likewise project them onto a 2-D disc. We generate each such vector as follows:

- 1) Sample \mathbf{z}_i ($i = 1, \dots, n-2$) from a standard normal distribution.
- 2) For $i = 1, \dots, n-2$, set $\hat{\mathbf{u}}_i = \frac{\sqrt{1-q^2} \times \mathbf{z}_i}{\sqrt{\sum_{i'=1}^{n-2} \mathbf{z}_{i'}^2}}$.

We then project the vectors in the set $\{\hat{\mathbf{u}}^{(j)}\}$ onto the two principal axes obtained from PCA. To enable a fair comparison between the variances of the randomly generated elements of \mathcal{T}^n and those obtained from the trained detectors, we run PCA *separately* for each set.

E. Results

We projected all vectors $\{\hat{\mathbf{u}}^{(j)}\}$ whose correlation with \mathbf{u} was between 0.575 and 0.625; this amounted to 55% of




Fig. 4. Coverage of the \mathcal{T}^n sphere (projected onto 2 dimensions using PCA) from different neural networks trained to predict student engagement (top) and age (bottom) from face images. We used either convolutional (CNN) or feed-forward (MLP) neural networks. For comparison, we also sampled random vectors using the procedure from Section V-D.

the engagement detectors and 28% of the age detectors. The projections are shown for each task in Figure 4. First, we observe that there is some “spread” – the measurement vectors occupy different clusters on the sphere. This indicates that the same training data can still yield automatic measurements \hat{u} on testing data whose correlations with each other is far less than 1. In fact, for engagement recognition, the minimum correlation, over all pairs $(\hat{u}^{(j)}, \hat{u}^{(j')})$, was 0.56.

For engagement recognition, the MLP-based measurements and the CNN-based measurements each resided within their own clusters on the sphere, and these clusters did not overlap. This suggests that, even though both architectures yielded similar overall accuracies, they are making different kinds of estimation errors on the test set.

Finally, a quantitative comparison of the variance between the automatic measurements $\hat{u}^{(1)}, \dots, \hat{u}^{(m)}$ and random samples from \mathcal{T}^n indicates that varying the training configuration (architecture, hyperparameters) provides only limited ability

to cover the sphere: the variance in the vectors, as quantified as the sum of the trace of their covariance matrix, was statistically significantly less compared to randomly sampled points on \mathcal{T}^n ($p < 0.01$, 1-tailed, Monte Carlo simulation).

VI. CONCLUSIONS

Advances in automatic face and recognition present a powerful opportunity to create new scientific instruments that can benefit basic research in sociobehavioral sciences. However, since detectors are often trained to estimate existing measures, which are already only an estimate of underlying constructs, then these instruments are essentially one step further removed from ground-truth. For this reason, it is important to interpret results obtained with them with care.

In this paper, we investigated how measurements of construct U obtained with an automatic detector can impact the estimated correlation between U and another construct V . We showed that: (1) The set of 0-mean unit-length n -vectors with a fixed Pearson product-moment correlation q to vector \mathbf{u} is a $(n - 3)$ -sphere \mathcal{T}^n embedded in \mathbb{R}^n . (2) If the correlation between automatic measurements \hat{u} and the ground-truth measurements is q ; if the true correlation between U and V is r ; and if \hat{u} is sampled uniformly from \mathcal{T}^n ; then the expected sample correlation obtained with the automatic detector is qr . (3) The probability of a “false correlation”, i.e., a sample correlation between constructs U and V whose sign differs from the true correlation, is monotonically decreasing in n (number of participants) and also monotonically decreasing in q (accuracy of the detector). These probabilities can be non-trivial for small values of n that are nonetheless sometimes found in contemporary research using automatic facial expression and affect detectors. Moreover, the danger of a false correlation is not eliminated through statistical significance testing. (4) We explored empirically how efficiently multiple neural network-based detectors of age and student engagement, when trained using different architectures and hyperparameters but the same training data, can “cover” the sphere \mathcal{T}^n .

In practice, our results suggest that, particularly when the number of participants is small and/or the accuracy of the detector is modest, it is important to consider the possibility of a false correlation, or at least a skewed correlation (by factor q), when drawing scientific conclusions.

Limitation and future work: In our study we assumed that \hat{u} is a random sample from the uniform distribution over \mathcal{T}^n – this expresses the idea that *a priori* we may have no idea which *particular* element of \mathcal{T}^n detector d will return. In reality, however, detectors have biases – e.g., due to head pose, lighting conditions, training set composition, etc. – and these can affect which element of \mathcal{T}^n is obtained.

VII. APPENDIX

A. Proof of Proposition 2

We prove the proposition for the case that $r < 0$; the case for $r > 0$ is similar.

From Section III, we have that

$$\rho(\hat{\mathbf{u}}, \mathbf{v}) = qr + \hat{u}_3 \sqrt{1 - r^2}$$

Since each \hat{u}_i ($i = 3, 4, \dots, n$) is a coordinate on an $(n-3)$ -sphere, it can be re-parameterized [12] by sampling $n-2$ standard normal random variables and normalizing, i.e.:

$$\hat{u}_i = \frac{\sqrt{1-q^2} \times z_i}{\sqrt{\sum_{j=3}^n z_j^2}}$$

where each $z_i \sim \mathbb{N}(0, 1)$. A false positive correlation thus occurs when $\rho(\hat{\mathbf{u}}, \mathbf{v})$ is at least $c = |qr|/\sqrt{1-r^2}$ more than its expected value qr :

$$\Pr[\hat{u}_3 \geq c] = \Pr \left[\frac{\sqrt{1-q^2} \times z_3}{\sqrt{\sum_{j=3}^n z_j^2}} \geq c \right]$$

Due to the inequality, we must handle the cases that $z_3 \geq 0$ and $z_3 < 0$ separately. Note that the latter case contributes 0 probability since $c \geq 0$ and $q > 0$. Also, since z_3 is a standard normal random variable, $\Pr[z_3 \geq 0] = 0.5$.

$$\begin{aligned} \Pr[\hat{u}_3 \geq c] &= \Pr \left[\frac{\sqrt{1-q^2} \times z_3}{\sqrt{\sum_{j=3}^n z_j^2}} \geq c \mid z_3 \geq 0 \right] \Pr[z_3 \geq 0] + \\ &\Pr \left[\frac{\sqrt{1-q^2} \times z_3}{\sqrt{\sum_{j=3}^n z_j^2}} \geq c \mid z_3 < 0 \right] \Pr[z_3 < 0] + \\ &= \frac{1}{2} \Pr \left[\frac{\sqrt{1-q^2} \times z_3}{\sqrt{\sum_{j=3}^n z_j^2}} \geq c \mid z_3 \geq 0 \right] + 0 \\ &= \frac{1}{2} \Pr \left[(1-q^2)z_3^2 \geq c^2 \sum_{j=3}^n z_j^2 \right] \\ &= \frac{1}{2} \Pr \left[(1-q^2-c^2)z_3^2 \geq c^2 \sum_{j=4}^n z_j^2 \right] \\ &= \frac{1}{2} \Pr \left[z_3^2 \geq \frac{c^2}{(1-q^2-c^2)} \sum_{j=4}^n z_j^2 \right] \end{aligned}$$

For $n > 3$, each side of the inequality is a sum of squared normally distributed random variables, i.e., a χ^2 -random variable (though with different degrees of freedom). We can thus rewrite this probability as

$$\begin{aligned} \Pr[\hat{u}_3 \geq c] &= \frac{1}{2} \Pr \left[\chi_1^2 \geq \left(\frac{c^2}{1-q^2-c^2} \right) \chi_{(n-3)}^2 \right] \\ &= \frac{1}{2} \int_0^\infty f_1(t) F_{n-3} \left(\frac{1-q^2-c^2}{c^2} t \right) dt \\ &\doteq h(n, q, r) \end{aligned}$$

where χ_1^2 and $\chi_{(n-3)}^2$ are χ^2 random variables with 1 and $(n-3)$ degrees of freedom, respectively. The probability is equivalent to the integral because, for any value t of the χ_1^2 variable, we require that the χ_{n-3}^2 variable be less than t (after applying a scaling factor). To our knowledge, there is

no closed formula for this integral, but we can compute it numerically. For $n = 3$, we have

$$\begin{aligned} \Pr[\hat{u}_3 \geq c] &= \frac{1}{2} \Pr[(1-q^2-c^2)z_3^2 \geq 0] \\ &= \frac{1}{2} \Pr[c^2 \leq 1-q^2] \end{aligned}$$

since a χ^2 -random variable is non-negative, and where the probability of $c^2 \leq 1-q^2$ is 1 if the inequality is true and 0 otherwise.

B. Proof of Proposition 3

For convenience, define $\alpha = \frac{1-q^2-c^2}{c^2}$.

$$\begin{aligned} h(n+1, q, r) - h(n, q, r) &= \int_0^\infty f_1(t) F_{(n+1)-3}(\alpha t) dt - \\ &\int_0^\infty f_1(t) F_{n-3}(\alpha t) dt \\ &= \int_0^\infty f_1(t) [F_{n-2}(\alpha t) - F_{n-3}(\alpha t)] dt \end{aligned}$$

Ghosh [5] proved that, for any fixed $t > 0$, $\Pr[\chi_k^2 > t]$ is monotonically increasing in the degrees of freedom k ; hence, $F_k(t)$ is monotonically decreasing in k . Therefore,

$$F_{n-2}(\alpha t) - F_{n-3}(\alpha t) < 0$$

for all t . Since f_k is a non-negative function for all k , then the integral in Equation 3 must be negative; hence, h is monotonically decreasing in n for every $c > 0$ and $q \in (0, 1]$.

C. Proof of Proposition 4

First, we show that α is monotonically decreasing in q^2 :

$$\begin{aligned} \alpha(q) &= \frac{1-q^2-c^2}{c^2} \\ &= \frac{1-q^2-q^2r^2/(1-r^2)}{q^2r^2/(1-r^2)} \\ &= \frac{(1-r^2)(1-q^2)-q^2r^2}{q^2r^2} \\ &= \frac{1-r^2-q^2}{q^2r^2} \\ &= \frac{1-r^2}{q^2r^2} - \frac{1}{r^2} \end{aligned}$$

The first term is monotonically decreasing in q^2 , and the second term is constant in q^2 .

Next, let ϵ be a positive real number such that $q + \epsilon \leq 1$:

$$\begin{aligned} h(n, q + \epsilon, r) - h(n, q, r) &= \int_0^\infty f_1(t) F_{n-3}(\alpha(q + \epsilon)t) dt - \\ &\int_0^\infty f_1(t) F_{n-3}(\alpha(q)t) dt \\ &= \int_0^\infty f_1(t) [F_{n-3}(\alpha(q + \epsilon)t) - F_{n-3}(\alpha(q)t)] dt \end{aligned}$$

Since F_{n-3} is monotonically *increasing*, then the expression in brackets is negative. Since f_1 is non-negative, then the entire integral must be less than 0.

REFERENCES

- [1] K. D. Carlson and A. O. Herdman. Understanding the impact of convergent validity on research results. *Organizational Research Methods*, 15(1):17–32, 2012.
- [2] L. Chen, S.-Y. Yoon, C. W. Leong, M. Martin, and M. Ma. An initial analysis of structured video interviews by using multimodal emotion detection. In *Proceedings of the 2014 workshop on Emotion Representation and Modelling in Human-Computer-Interaction-Systems*, pages 1–6. ACM, 2014.
- [3] N. L. Collins and S. J. Read. Adult attachment, working models, and relationship quality in dating couples. *Journal of personality and social psychology*, 58(4):644, 1990.
- [4] A. L. Duckworth, C. Peterson, M. D. Matthews, and D. R. Kelly. Grit: perseverance and passion for long-term goals. *Journal of personality and social psychology*, 92(6):1087, 2007.
- [5] B. Ghosh. Some monotonicity theorems for χ^2 , F and t distributions with applications. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 480–492, 1973.
- [6] G. V. Glass and J. R. Collins. Geometric proof of the restriction on the possible values of r_{xy} when r_{xz} and r_{yz} are fixed. *Educational and Psychological Measurement*, 30(1):37–39, 1970.
- [7] G. Huang, Y. Li, G. Pleiss, Z. Liu, J. E. Hopcroft, and K. Q. Weinberger. Snapshot ensembles: Train 1, get m for free. *arXiv preprint arXiv:1704.00109*, 2017.
- [8] M. P. Lab. The MPLab GENKI Database. <http://mplab.ucsd.edu>.
- [9] C.-K. Leung and K. Lam. A note on the geometric representation of the correlation coefficients. *The American Statistician*, 29(3):128–130, 1975.
- [10] G. Matthews, L. Joyner, K. Gilliland, S. Campbell, S. Falconer, and J. Huggins. Validation of a comprehensive stress state questionnaire: Towards a state big three. *Personality psychology in Europe*, 7:335–350, 1999.
- [11] H. Monkaresi, N. Bosch, R. A. Calvo, and S. K. D’Mello. Automated detection of engagement using video-based estimation of facial expressions and heart rate. *IEEE Transactions on Affective Computing*, 8(1):15–28, 2017.
- [12] M. E. Muller. A note on a method for generating points uniformly on n -dimensional spheres. *Communications of the ACM*, 2(4):19–20, 1959.
- [13] I. Olkin. Range restrictions for product-moment correlation matrices. *Psychometrika*, 46(4):469–472, 1981.
- [14] F. Parra, R. Miljkovitch, G. Persiaux, M. Morales, and S. Scherer. The multimodal assessment of adult attachment security: developing the biometric attachment test. *Journal of medical Internet research*, 19(4), 2017.
- [15] G. Perugia, D. Rodríguez-Martín, M. D. Boladeras, A. C. Mallofré, E. Barakova, and M. Rauterberg. Electrodermal activity: explorations in the psychophysiology of engagement with social robots in dementia. In *Robot and Human Interactive Communication (RO-MAN), 2017 26th IEEE International Symposium on*, pages 1248–1254. IEEE, 2017.
- [16] M.-Z. Poh, N. C. Swenson, and R. W. Picard. A wearable sensor for unobtrusive, long-term assessment of electrodermal activity. *IEEE transactions on Biomedical engineering*, 57(5):1243–1252, 2010.
- [17] H. F. Priest. Range of correlation coefficients. *Psychological reports*, 22(1):168–170, 1968.
- [18] L. N. Smith. Cyclical learning rates for training neural networks. In *Applications of Computer Vision (WACV), 2017 IEEE Winter Conference on*, pages 464–472. IEEE, 2017.
- [19] J. Whitehill, Z. Serpell, Y.-C. Lin, A. Foster, and J. R. Movellan. The faces of engagement: Automatic recognition of student engagement from facial expressions. *IEEE Transactions on Affective Computing*, 5(1):86–98, 2014.